



CCs

numberfy

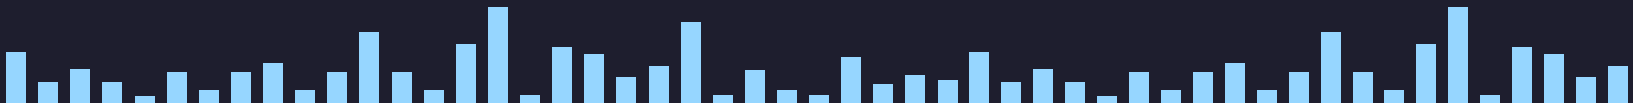
→ whoami

Joaquin Badillo

Andres Tarazona

Rodrigo Núñez

Alejandro Arouesty






Agenda


→ Is agenda/

Introduction	1.1
Cloud Architecture	1.2
Auth	1.3
Databases	1.4
Load Balancing & Scalability	1.5
UNet AutoEncoder	1.6
Demo	1.7
Q&A	1.8
[Appendix]	

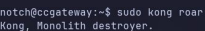
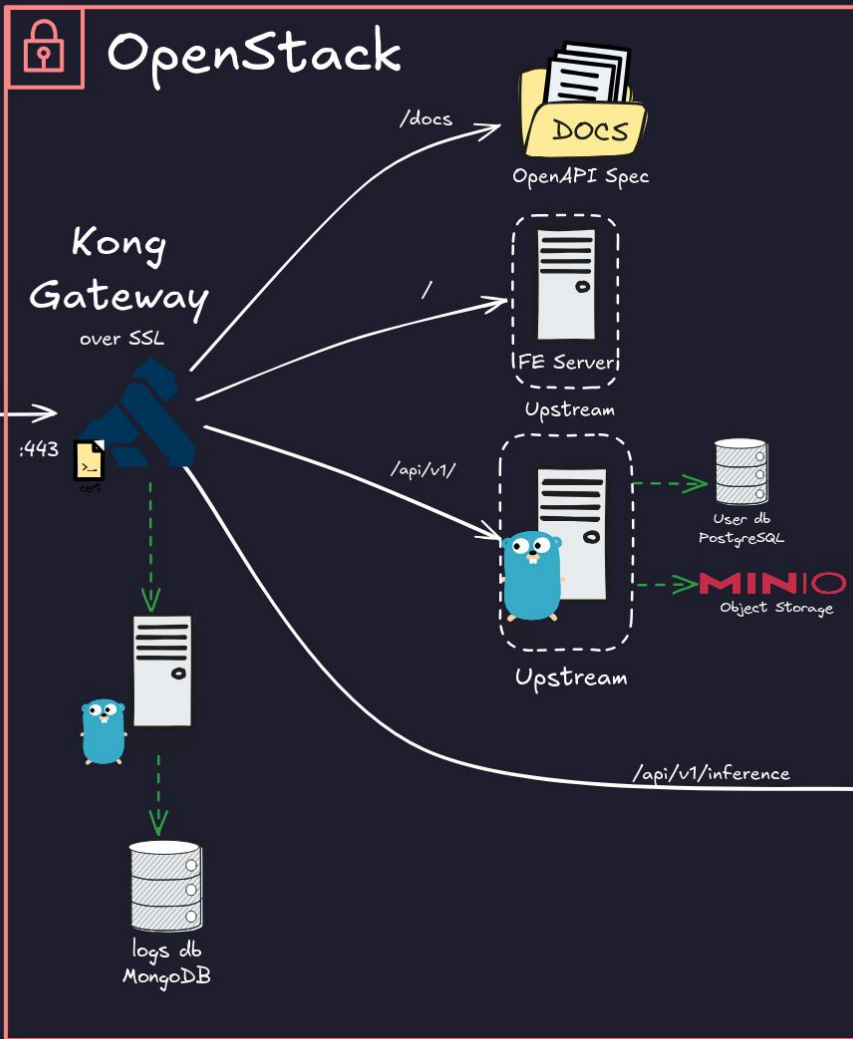




Introduction


- Web application that allows drawing digits (0-9) on a 112x112 canvas and getting them recognized.
 - Automatic recognition system using AI.
 - Inference history per user.
 - Authentication flows with OAuth2 standard implementation.
- 

Year	Percentage
2008	25
2009	28
2010	42
2011	42
2012	22
2013	35
2014	52
2015	48
2016	45
2017	55

[illegible]




Load Balancing and Scalability

- Centralized API Gateway
 - 2 load-balanced Frontend instances
 - 2 load-balanced Backend instances
 - Centralized logging on dedicated VM
 - PAT Overload for external access from Tec network
- 



Authentication

- Framework: Gin + GORM with PostgreSQL
 - OAuth2 Flows: Authorization Code, Client Credentials, Password Grant, Refresh Token
 - Key Endpoints:
 - `/oauth2/authorize` → Authorization code generation
 - `/oauth2/token` → Token exchange (access + refresh)
 - `/oauth2/introspect` → Token validation
 - Authentication Middleware: Bearer token validation for protected routes.
 - Administration: Complete management of clients, consumers and tokens.
- 



Databases

Robust PostgreSQL Schema

- Users: Authentication with bcrypt, UUIDs as PKs
- Images: Tracking of pairs (sent_image_id, received_image_id)
- OAuth2: Credentials, Tokens, Authorization Codes, Refresh Tokens

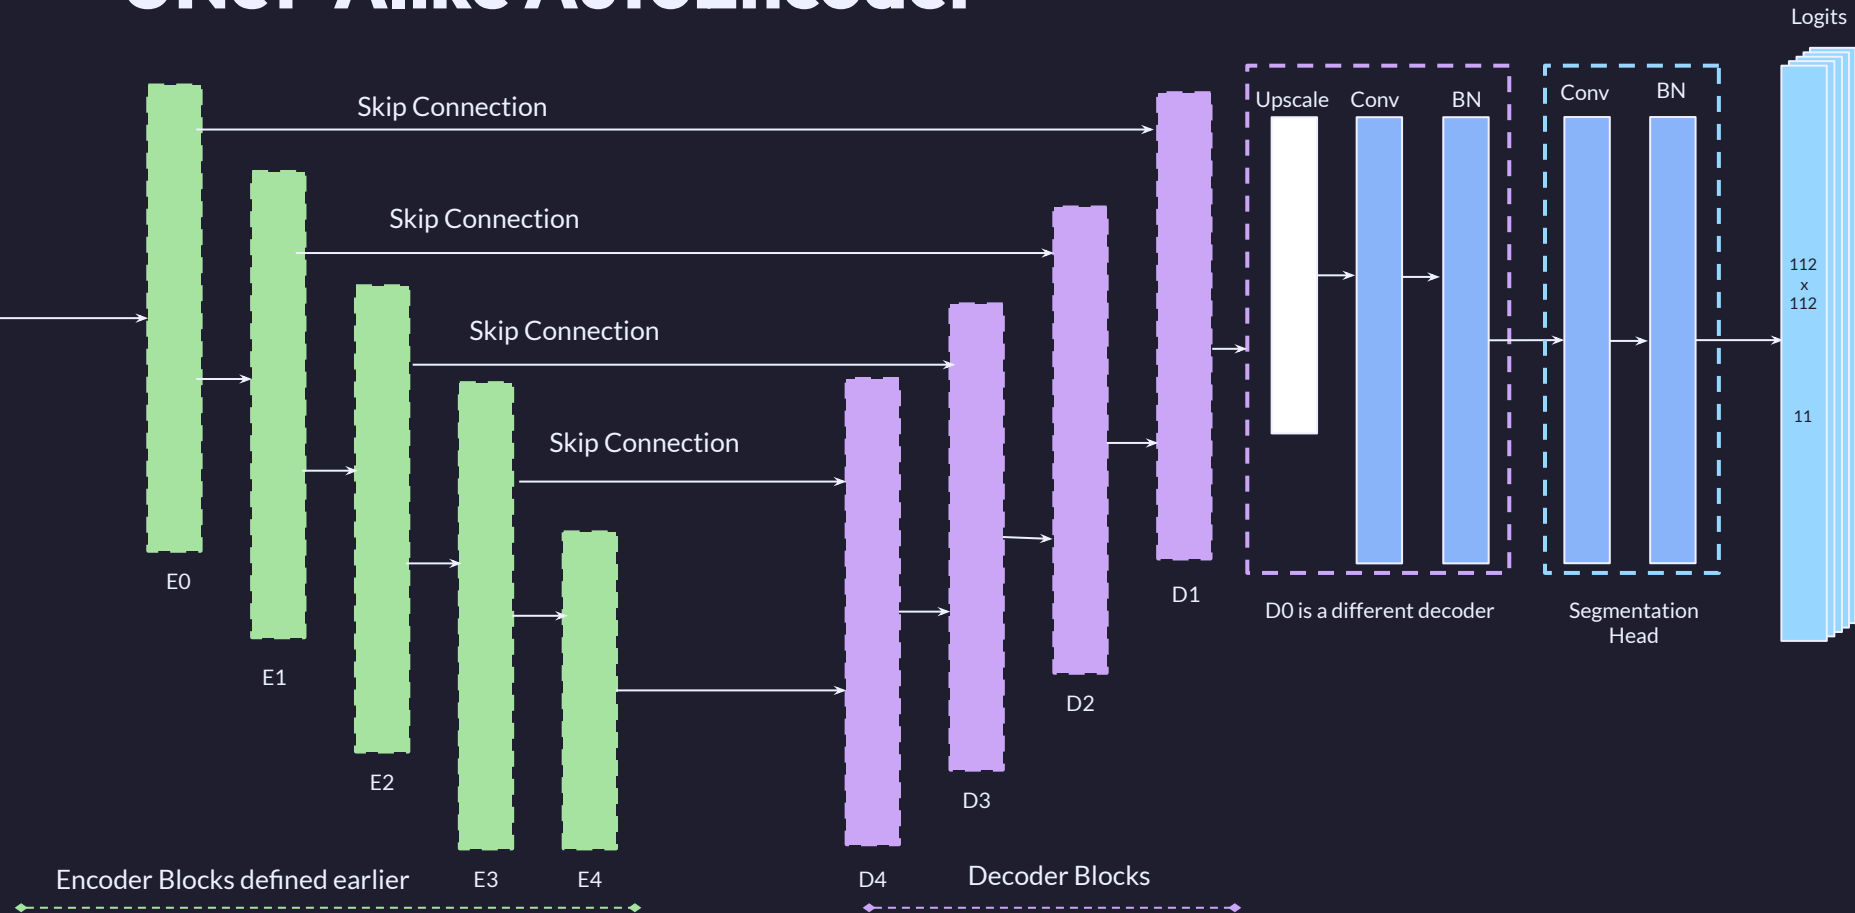
Object Storage using MinIO (buckets)

- Image IDs are stored in PostgreSQL

NoSQL MongoDB instance for Logs



UNet-Alike AutoEncoder



Check the Notebooks




<https://drive.google.com/drive/folders/1vhfsJTQdxCH4m2n7vYLj330JxrSGqNQg?usp=sharing>



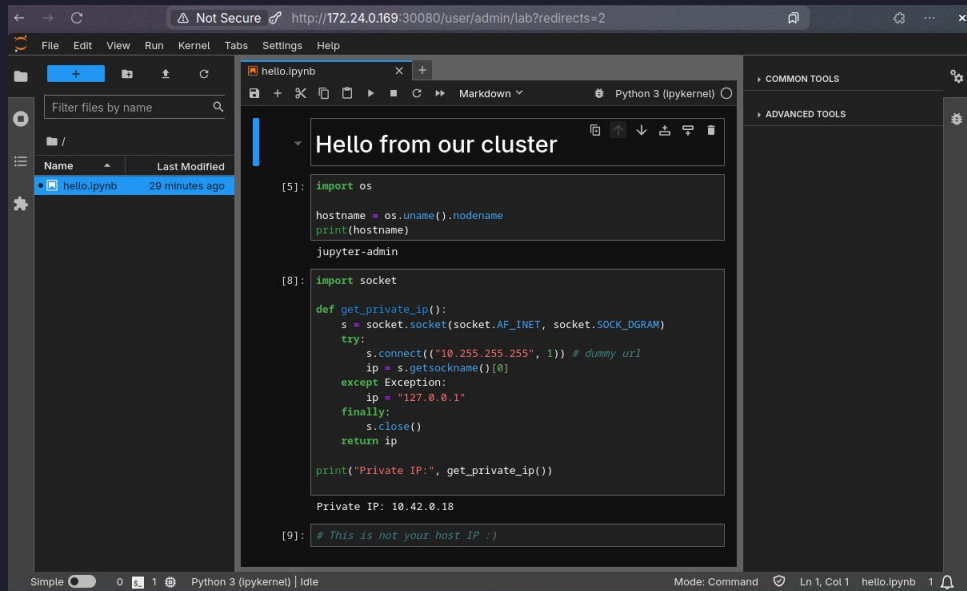
DEMO

Click me ↑



→ curl <https://10.49.12.47:8443/>

Future Work



The screenshot shows a JupyterLab environment. On the left is a file explorer with a table of files:

Name	Last Modified
hello.ipynb	29 minutes ago

The main area displays a Python notebook titled 'hello.ipynb' with the following code:

```
[5]: import os
hostname = os.uname().nodename
print(hostname)
jupyter-admin

[8]: import socket

def get_private_ip():
    s = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
    try:
        s.connect(("10.255.255.255", 1)) # dummy url
        ip = s.getsockname()[0]
    except Exception:
        ip = "127.0.0.1"
    finally:
        s.close()
    return ip

print("Private IP:", get_private_ip())

Private IP: 10.42.0.18

[9]: # This is not your host IP :)
```

→ echo “MLOps”




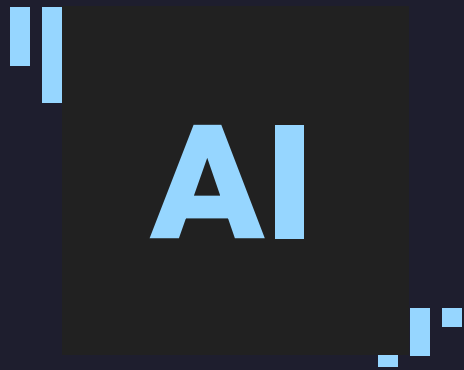
→ echo “Authorization codes”
&&
echo “Further Scopes”



Q&A

→ echo “Thank you”

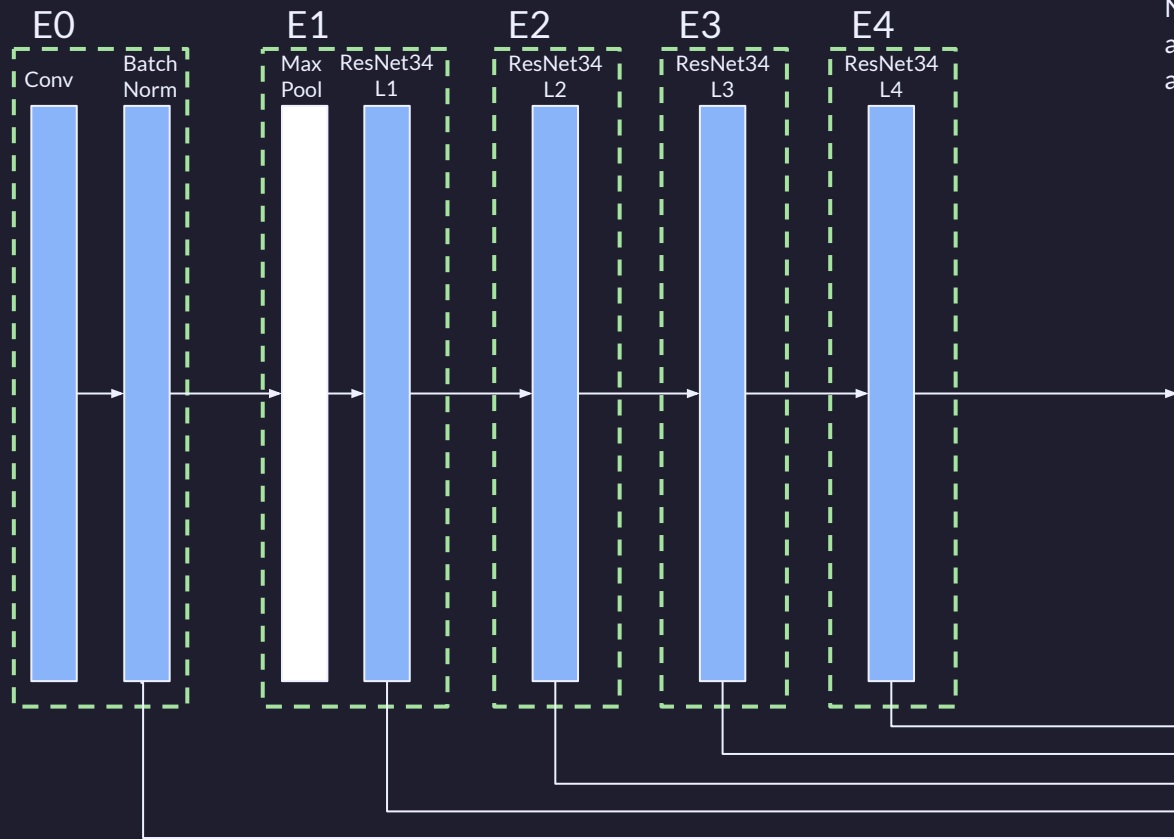




Appendix: UNet Architecture

> An idiot admires complexity, a genius admires simplicity
- Terry Davis

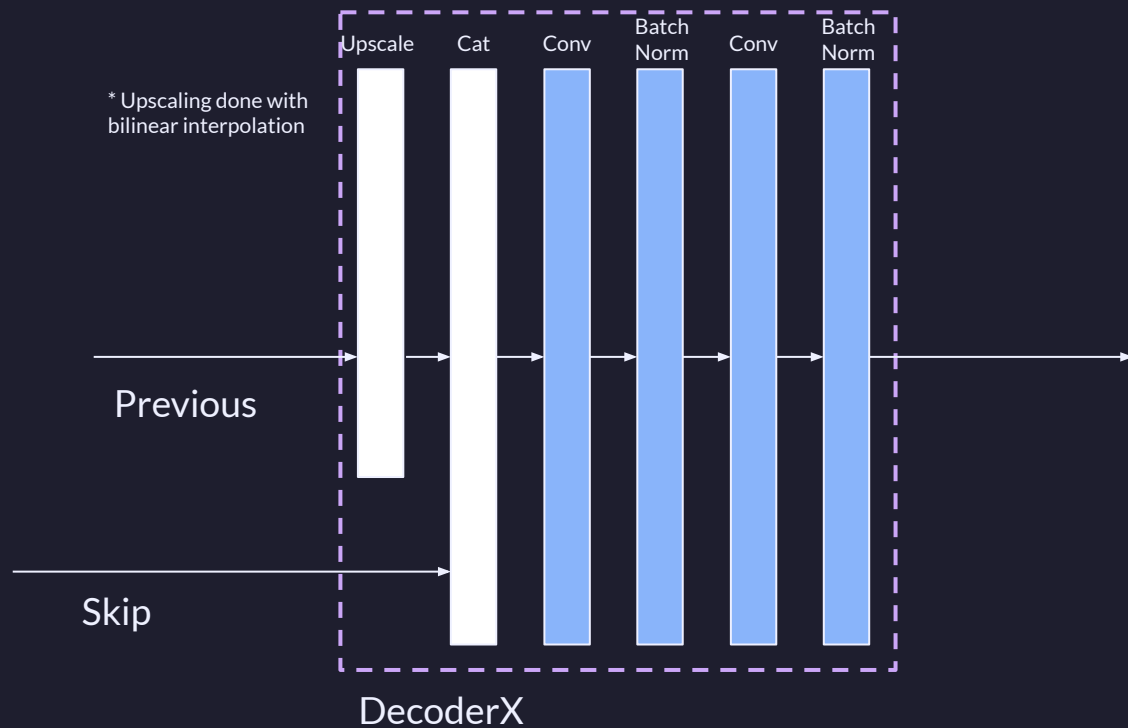
Encoder Blocks



Note that ResNet Layers are abstracted away (some encoder layers of ResNet also have skip connections)

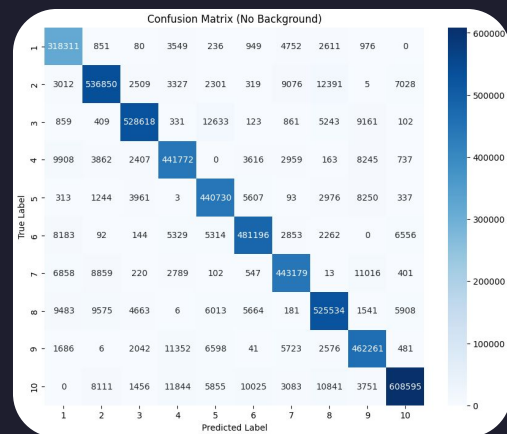


Decoder Block



Results: Model Performance

	Prec	Recall	F1	Acc
1	0.888	0.958	0.921	0.958
2	0.942	0.931	0.936	0.931
3	0.968	0.947	0.957	0.947
4	0.920	0.933	0.926	0.933
5	0.919	0.951	0.934	0.951
6	0.947	0.940	0.943	0.940
7	0.937	0.935	0.936	0.935
8	0.931	0.924	0.927	0.924
9	0.915	0.938	0.926	0.938
0	0.966	0.917	0.941	0.917



Results

